PROJECT 01 · CURRENT

EduLoans · CashKaro

B2C education-loan workflow · 8 lender partners · multi-role

62	23	686	5
STRUCTURED BRE RULES	WORKFLOW STATES	LoC LIVE BRE FN	AI TOOLS CHAINED

01 THE PROBLEM

What was broken?

Education loan underwriting in India is opaque. Each lender has its own rules, its own document checklist, its own appetite for borderline cases. Students see a confusing market. Distribution partners can't track lead status. Key Account Managers (KAMs) firefight stalled cases manually. The product had to fix all three at once.

02 THE DIAGNOSIS

What signals confirmed the shape of it?

- 8 lender partners · 4 applicant types = **32 distinct rule combinations** needed for the document checklist alone — a flat master list would over-collect and under-collect simultaneously.
- Lead-source data showed **Punjab pincodes and certain channels behaving as junk** — needed geo-exclusion at intake, not after the engine ran.
- Indian-document reality (regional marksheets, scanned ITRs, joint bank statements) **breaks pure-OCR pipelines** — designing the human-in-loop step before launch costs less than retrofitting after the first 1,000 stuck cases.

03 THE DECISION

Deterministic engine over ML scoring on day one.

Deterministic = auditable, defensible to lender RMs, explainable to ops. ML scoring sits in v2, once we have outcome data. Documented this trade-off in the PRD so it isn't relitigated every sprint.

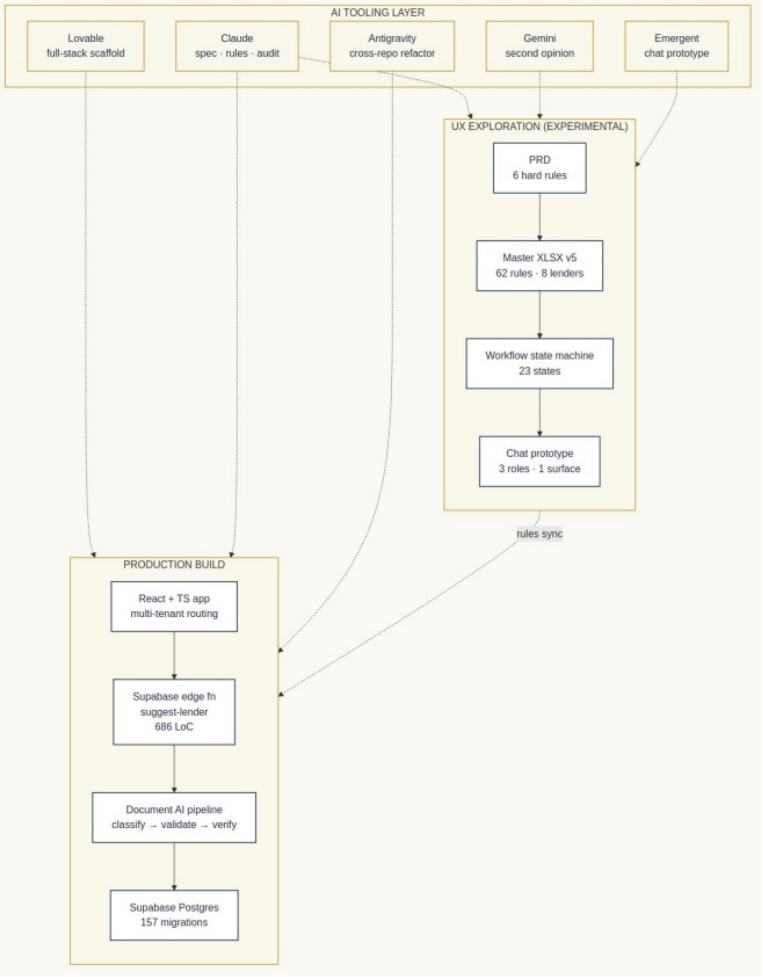
Three roles, one chat surface — not three dashboards.

Conventional fintech instinct is three dashboards. I bet on a single conversational surface with a role pill. Unified mental model > role-specific feature density at this stage.

KAM verification gates BRE display.

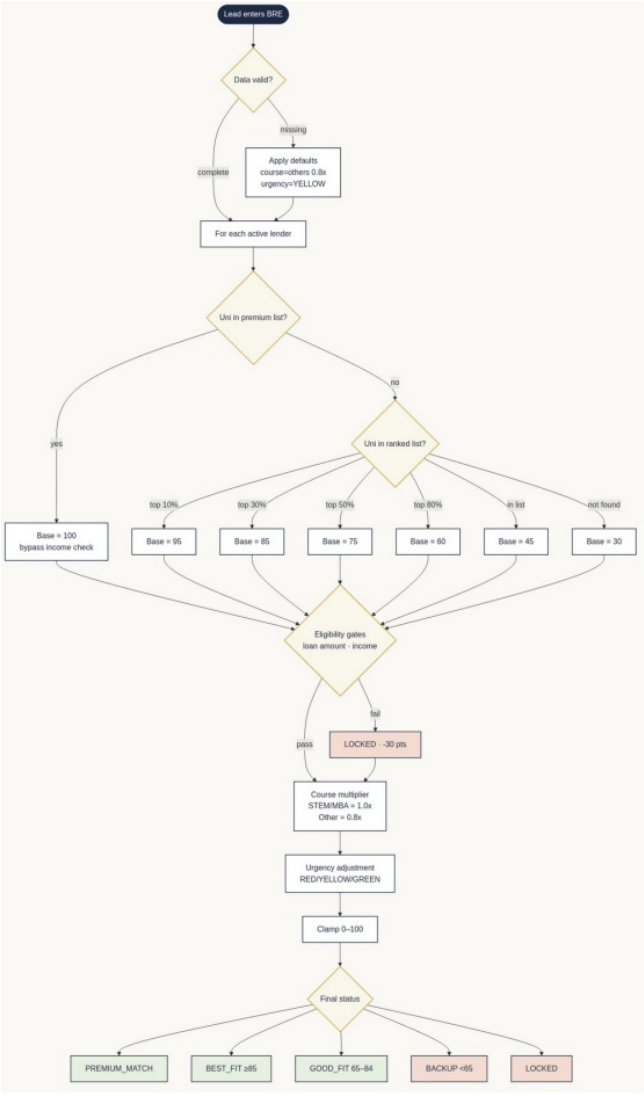
Showing a BRE result to an unverified lead anchors expectations and erodes lender trust. Enforced at the workflow-state-machine level, not just the UI.

Architecture



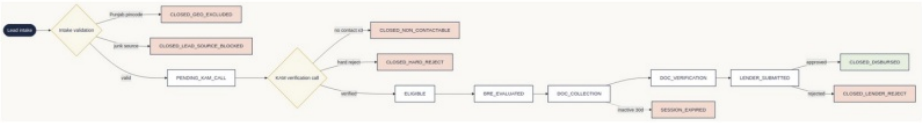
System architecture · AI tooling layer above, deterministic core below

How the BRE decides



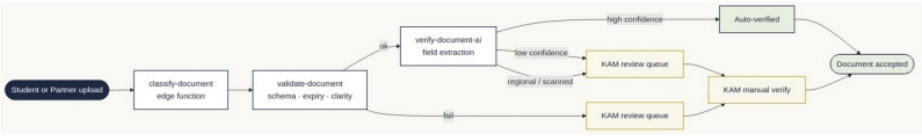
Per-lender deterministic scoring · same input, same output, every time

Lead workflow



23 canonical states · BRE never runs before KAM verification call

Document AI pipeline



classify → validate → verify · KAM review queues for low-confidence and regional documents

04 HOW IT EXECUTED

1	**Audit master XLSX** · Logged 10 contradictions, 12 missing values, 10 weak spots in the source-of-truth file
2	**Author PRD** · Six non-negotiable hard rules · scope cuts · v1 vs v2 boundary set explicitly
3	**Model 62 BRE rules** · Across 8 lenders + 5 global, plus 3 lookup tabs (lender, university rank, pincode tier)
4	**Build production scaffold** · Multi-tenant React + Supabase via Lovable; 686-LoC deterministic edge function
5	**Build experimental UX** · Chat-first three-role prototype to test unified-surface hypothesis (see Project 02)
6	**Document & sign off** · Workflow state machine, KAM dashboard wireframe, all 15 open questions resolved with product

05 AI TOOLING STACK

LOVABLE	Full-stack scaffold of multi-tenant React + Supabase app, auth, routes, BRE edge function. Fastest path from prompt to deployed full-stack app.
EMERGENT	Chat-first prototype scaffold for the three-role conversational UI. Right tool when the product is conversation-shaped.
CLAUDE	Reasoning layer — PRD, BRE rule modelling, conflict-and-gap audit on master XLSX, code review across the production stack.
GEMINI	BRE scoring tradeoffs, edge-case enumeration, second opinion on rule design. Different failure modes than Claude — useful for debiasing.
ANTIGRAVITY	Cross-repo orchestration, bulk refactor, TypeScript debt cleanup, integration sweeps. Agentic IDE that holds context across the codebase.

06 WHAT I LEARNED

Modern AI tools are not interchangeable.

Pick each one for the specific job it does best, then make sure the handoffs are clean. Lovable for scaffold, Claude for reasoning, Antigravity for refactor. Each one earned its place by being a poor substitute for the others' strengths.

Deterministic core, agentic edges, human-in-the-loop where the cost of being wrong is real.

This is the durable architecture for regulated workflows. EduLoans is one example. The Absolute work was an earlier instance of the same pattern. The job is identifying which boundary you're working on — and never letting the tool choose for you.

Repo · github.com/priyamsameer95-sys/eduloans

PROJECT 02 · EXPERIMENTAL

EduLoans · UX exploration

Three roles, one chat surface · mobile-first prototype

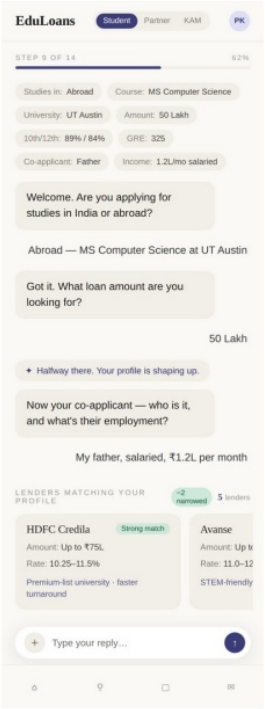
3 → 1	14	23	1
ROLES UNIFIED	TURN STUDENT FLOW	STATES MAPPED	WEEK TO PROTOTYPE

01 THE HYPOTHESIS

What was the bet?

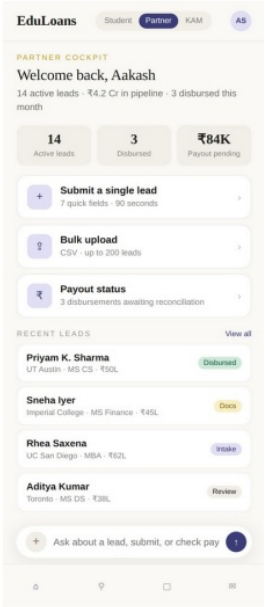
The main build splits Student, Partner, and KAM into separate dashboards. I wanted to test whether a single conversational surface — same shell, role-specific seeded conversations and action sheets — would feel more natural than three different products stitched together.

02 THREE ROLES, ONE SHELL



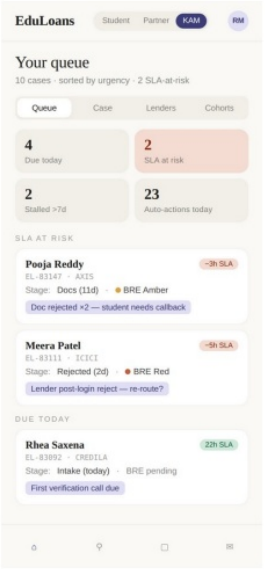
Student

14-turn capture · live lender narrowing



Partner

Submit · bulk upload · payouts



KAM

Queue · SLA risk · case detail

03 WHAT I LEARNED

Unified surface beats role-specific feature density at this stage.

All three roles share the same primitives — a feed, a profile chip strip, an action pill, a queue view. Splitting them across three apps multiplies engineering, design and support cost without adding clarity. The prototype made the case visually.

Where it stays a prototype, not the production build.

The main build needs multi-tenant routing, auth, role-based dashboards for ops at scale. The chat surface is a candidate for the next iteration once the underlying engine is settled — not a replacement.

Repo · github.com/priyamsameer95-sys/V1.03-Chat-eduloans